

# Boosting: boosted trees and gradient boosting

Wei Li

Syracuse University

Spring 2024

# OVERVIEW

Boosting trees

Gradient boosting

Tuning hyperparameters

# Boosting trees

# Boosting trees

A CART tree can be expressed as

$$T(x; \Theta) = \sum_{j=1}^J \gamma_j 1(x \in R_j)$$

with parameters  $\Theta = (\gamma_j, R_j)_1^J$ .

$$\hat{\Theta} = \arg \min_{\Theta} \sum_{j=1}^J \sum_{x_i \in R_j} L(y_i, \gamma_j) = \arg \min_{\Theta} \sum_{i=1}^n L(y_i, T(x_i, \Theta))$$

- ▶ Finding  $\gamma_j$  given  $R_j$  : Given the  $R_j$ , estimating the  $\gamma_j$  is typically trivial, and often  $\hat{\gamma}_j = \bar{y}_j$ , the mean of the  $y_i$  falling in region  $R_j$ .
  - ▶ For classification,  $\hat{\gamma}_j$  is the majority class of the observations falling in region  $R_j$ .
- ▶ Finding  $R_j$  : A typical strategy is to use a **greedy, top-down recursive partitioning algorithm** to find the  $R_j$ .

## Two modifications

We solve the problem iteratively building the model.

We approximate the objective function by a smoother and more convenient criterion for optimizing the  $R_j$

$$\tilde{\Theta} = \arg \min_{\Theta} \sum^n \tilde{L}(y_i, T(x_i, \Theta))$$

Then given the  $\hat{R}_j = \tilde{R}_j$ , the  $\gamma_j$  can be estimated more precisely using the original criterion.

# A simple algorithm for boosting regression trees

Assume  $J_m = J$  fixed (same amount of terminal nodes).

1. Set  $\hat{f}(x) = 0$  and  $r_i = y_i$  for all  $i$  in the training set.
2. For  $m = 1, 2, \dots, M$ , repeat:
  - (a) Fit a tree  $\hat{f}^m$  with  $J$  terminal nodes to the training data  $(X, r)$
  - (b) Update  $\hat{f}$  by adding in a shrunk version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}_m(x)$$

- (c) Update the residuals,

$$r_i \leftarrow r_i - \lambda \hat{f}_m(x_i)$$

3. Output the boosted model,

$$\hat{f}(x) = \sum_{m=1}^M \lambda \hat{f}_m(x)$$

- ▶  $M$ : too large of  $M$  leads to overfit
- ▶  $\lambda$ : often choose some small  $\lambda = 0.01$  or  $0.001$ .
  - ▶ Smaller  $\lambda$  leads to larger  $M$  needed.
- ▶  $J$  controls the interaction depth (need not be too large).
  - ▶  $J = 1$  is a stump.

# Boosted trees

The boosted tree model is a sum of such trees,

$$f_M(x) = \sum_{m=1}^M T(x; \Theta_m)$$

induced in a forward stagewise manner.

At each step  $m$ , given the current tree model  $f_{m-1}(x)$ , we solve

$$\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^n L(y_i, f_{m-1}(x_i) + T(x_i, \Theta_m))$$

where  $\Theta_m = (\gamma_{jm}, R_{jm}), j = 1, \dots, J_m$  the parameter of the next tree.

$$\hat{\Theta}_m = (\hat{\gamma}_{jm}, \hat{R}_{jm}) = \arg \min_{\Theta_m} \sum_{i=1}^n L(y_i, f_{m-1}(x_i) + T(x_i, \Theta_m)) \quad (1)$$

- Find  $\gamma_{jm}$  given  $R_{jm}$  - easy

$$\hat{\gamma}_{jm} = \arg \min_{\gamma_{jm}} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma_{jm}) \quad (2)$$

- Find  $R_{jm}$  's - not so straightforward, except for squared-error loss and exponential loss (binary).
  - Squared-error loss. At each stage, the solution (1) to is a fitted regression tree to residuals  $y_i - f_{m-1}(x_i)$ ,  $\hat{\gamma}_{jm}$  is the mean of these residuals in each corresponding region.
  - exponential loss
    - Adaboost
    - real Adaboost



## Two-class classification: AdaBoost(M.1)

If all trees  $T(x; \Theta_m)$  are restricted to **scaled classification trees**, i.e., trees of the form  $\beta_m T(x; \Theta_m)$  with optimal constant  $\gamma_{jm}$  (on the region  $R_{jm}$ ) restricted to  $\{-1, 1\}$  (i.e. discrete), then solving 1

$$\arg \min_{\Theta_m, \beta_m} \sum_{i=1}^n \exp \left( -y_i \left( f_{m-1}(x_i) + \beta T(x_i, \Theta_m) \right) \right)$$
$$\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^n w_i^{(m)} 1(y_i \neq T(x_i; \Theta_m)) \quad (3)$$

- ▶ weights  $w_i^{(m)} = \exp(-y_i f_{m-1}(x_i))$
- ▶  $\beta_m = \log((1 - \text{err}_m)/\text{err}_m)/2$

Fit a small tree  $\in \{-1, 1\}$  minimizing the above *weighted 0/1 loss* (3)

# Real AdaBoost

**Real AdaBoost:** impose no restriction on the type of tree (same exponential loss, with no  $\beta$ ), the problem (1) is

$$\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^n w_i^{(m)} \exp(-y_i T(x_i; \Theta_m))$$

The solution to (2) given  $R_{jm}$  is

$$\hat{\gamma}_{jm} = \frac{1}{2} \log \frac{\sum_{x_i \in R_{jm}} w_i^{(m)} 1(y_i = 1)}{\sum_{x_i \in R_{jm}} w_i^{(m)} 1(y_i = -1)}$$

with weights  $w_i^{(m)} = \exp(-y_i f_{m-1}(x_i))$ .

This requires a specialized tree-growing algorithm.

# Gradient boosting

# Gradient boosting

Consider the numerical optimization

$$\arg \min_{f \in \{f_M\}} L(f) \equiv \sum_{i=1}^n L(y_i, f(x_i))$$

The solution  $f_M$  is written as a sum of component vectors

$$\mathbf{f}_M = \sum_{m=0}^M \mathbf{h}_m, \quad h_m \in \mathbb{R}^n$$

where  $\mathbf{f}_M = (f_M(x_1), \dots, f_M(x_n))^{\top}$  and  $\mathbf{h}_m$  is the increment vector at the  $m$  th step.

The general idea of a greedy algorithm is to solve:

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$$
$$f_m(x) = f_{m-1}(x) + \arg \min_{h_m} [\sum_{i=1}^n L(y_i, f_{m-1}(x_i) + h_m(x_i))]$$

where  $h_m$  is a base learner function.

Unfortunately, choosing the best function  $h$  at each step for an arbitrary loss function  $L$  is a computationally infeasible optimization problem in general.

We consider **functional gradient descent** method.

# functional gradient descent

- ▶ Steepest descent: choose  $\mathbf{h}_m = -\rho_m \mathbf{g}_m$ , where  $\rho_m$  is a scalar and the  $i$ th component of  $\mathbf{g}_m$  is given by

$$g_{im} := g_m(x_i) = \left[ \frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f(x_i)=f_{m-1}(x_i)}$$

- ▶ The step length  $\rho_m$  is the solution to

$$\rho_m = \arg \min_{\rho} \sum_{i=1}^n L(y_i, f_{m-1}(x_i) - \rho g_m(x_i))$$

- ▶ Solution update:  $f_m(x_i) = f_{m-1}(x_i) - \rho_m g_m(x_i)$

# A comparison

$$\rho_m = \arg \min_{\rho} \sum_{i=1}^n L(y_i, f_{m-1}(x_i) - \rho g_m(x_i))$$

compared with Forward stagewise boosting trees:

$$\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^n L(y_i, f_{m-1}(x_i) + T(x_i; \Theta_m))$$

- ▶ Tree predictions  $T(x_i; \Theta_m)$  are analogous to the negative gradients  $-g_m(x_i)$
- ▶ But  $\mathbf{t}_m = \{T(x_1; \Theta_m), \dots, T(x_n; \Theta_m)\}$  are not independent, but constrained to be predictions of a  $J_m$ -terminal node decision tree, whereas  $-g_m$  is the unconstrained maximal descent direction.
- ▶  $\rho_m = \arg \min_{\rho} L(y, f_{m-1} - \rho g_m)$  is analogous to  $\hat{\gamma}_{jm} = \arg \min_{\gamma_{jm}} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma_{jm})$ , but the latter is solved in each region.

The gradient is defined only at the training data points  $x_i$  whereas the ultimate goal is to generalize  $f_M(x)$  to new data not represented in the training set.

A possible solution is to fit the tree  $T$  to the negative gradient values by *least squares*:

- Fit a tree  $T(x; \Theta_m)$  at  $m$  th iteration whose predictions  $\mathbf{t}_m$  are as close as possible to the negative gradient

$$\tilde{\Theta}_m = \arg \min_{\Theta} \sum_{i=1}^n (-g_{im} - T(x_i; \Theta))^2$$

- From the (approximate) solution regions  $\tilde{R}_{jm}$  set,

$$\hat{\gamma}_{jm} = \arg \min_{\gamma_{jm}} \sum_{x_i \in \tilde{R}_{jm}} L(y_i, f_{m-1}(x_i) + \gamma_{jm})$$



# Gradient tree boosting

## Algorithm 10.3 (Gradient tree boosting)

1. Initialize  $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$
2. For  $m = 1$  to  $M$  :
  - (a) For  $i = 1, 2, \dots, n$  compute *pseudo-residuals*

$$r_{im} = - \left[ \frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}$$

- (b) Fit a regression tree to the targets  $r_{im}$  giving terminal regions

$$R_{jm}, j = 1, 2, \dots, J_m$$

- (c) For  $j = 1, 2, \dots, J_m$  compute

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma)$$

- (d) Update  $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} 1(x \in R_{jm})$

3. Output  $\hat{f}(x) = f_M(x)$ .

The following table summarizes the gradients for some loss functions.

Setting	Loss Function	$-\partial L(y_i, f(x_i))/\partial f(x_i)$
Regression	$\frac{1}{2}[y_i - f(x_i)]^2$	$y_i - f(x_i)$
Regression	$ y_i - f(x_i) $	$\text{sign}[y_i - f(x_i)]$
Regression	Huber	$y_i - f(x_i)$ for $ y_i - f(x_i)  \leq \delta_m$ $\delta_m \text{sign}[y_i - f(x_i)]$ for $ y_i - f(x_i)  > \delta_m$ where $\delta_m = \alpha\text{th-quantile}\{ y_i - f(x_i) \}$
Classification	Deviance	$k\text{th component: } I(y_i = \mathcal{G}_k) - p_k(x_i)$

ESL Table. 10.2. Gradients for commonly used loss functions.

# Gradient tree boosting for classification

For **K-class classification**, Class label  $Y \in \{1, \dots, K\}$ ,  $K \geq 3$ . Using one-hot encoding: let  $Y_i = (Y_{i1}, \dots, Y_{iK})$

►  $Y_{ik} = 1$  if  $i$ -th obs. in class  $k$ , 0 otherwise

For classification ( $K$ -classes),  $K$  separate trees are built at each iteration, producing  $p_k(x) = e^{f_k} / \sum_l e^{f_l}$ .

Recall assuming  $\sum_{k=1}^K f_k(x) = 0$ , multinomial deviance is

$$L(y, p(x)) = - \sum_{k=1}^K 1(y = k) f_k(x) + \log \left( \sum_{\ell=1}^K e^{f_{\ell}(x)} \right)$$

Each tree  $T_{km}$  is fit to its respective negative gradient vector  $g_{km}$ ,

$$\begin{aligned} -g_{ikm} &= \left[ \frac{\partial L(y_i, f_1(x_i), \dots, f_K(x_i))}{\partial f_k(x_i)} \right]_{\mathbf{f}(x_i) = \mathbf{f}_{m-1}(x_i)} \\ &= y_{ik} - p_k(x_i) \end{aligned}$$

**Algorithm 10.4** (Gradient tree boosting for classification with multinomial deviance loss)

1. Initialize  $f_{k0}(x) = 0, k = 1, 2, \dots, K$ .
2. For  $m = 1$  to  $M$  :
  - (a) Set

$$p_k(x) = \frac{e^{f_k(x)}}{\sum_{\ell=1}^K e^{f_\ell(x)}}, k = 1, 2, \dots, K$$

- (b) For  $k = 1$  to  $K$  :
  - i. Compute  $r_{ikm} = y_{ik} - p_k(x_i), i = 1, 2, \dots, n$
  - ii. Fit a regression tree to the targets  $r_{ikm}, i = 1, 2, \dots, n$  giving terminal regions  $R_{jkm}, j = 1, 2, \dots, J_m$
  - iii. Compute

$$\gamma_{jkm} = \frac{K-1}{K} \frac{\sum_{x_i \in R_{jkm}} r_{ikm}}{\sum_{x_i \in R_{jkm}} |r_{ikm}| (1 - |r_{ikm}|)}, j = 1, 2, \dots, J_m$$

- iv. Update  $f_{km}(x) = f_{k,m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jkm} 1(x \in R_{jkm})$
3. Output  $\hat{f}_k(x) = f_{kM}(x), k = 1, 2, \dots, K$ .

# Tuning hyperparameters

# Choosing tree size $J$

At each iteration a  $J$ -terminal node regression tree is induced. Thus  $J$  is the tuning parameter.

Other consideration

- ▶ The interaction level of tree-based approximations is limited by the tree size  $J$ .
- ▶ No interaction effects of level greater than  $J - 1$  are possible.
  - ▶ Setting  $J = 2$  (single split “decision stump”) produces boosted models with only main effects; no interactions are permitted.

Experience so far indicates that  $4 \leq J \leq 8$  works well in the context of boosting.

## Choosing $M$

### Shrinkage:

Line 2( $d$ ) of Algorithm 10.3 is replaced by

$$f_m(x) = f_{m-1}(x) + \nu \sum_{j=1}^J \gamma_{jm} 1(x \in R_{jm})$$

The parameter  $\nu$  is learning rate.

- ▶ smaller values of  $\nu$ , more shrinkage (less complexity)
- ▶ higher  $M$ , more complexity
- ▶ a tradeoff between them  $\nu$  and  $M$

Empirically the best strategy is to set  $\nu$  to be very small ( $\nu < 0.1$ ) and choose  $M$  by early stopping.

### Subsampling:

With stochastic gradient boosting (Friedman, 1999), at each iteration we sample a fraction of the training observations (without replacement), and grow the next tree using that subsample. The rest of the algorithm is identical.