

Boosting: boosted trees and gradient boosting

Wei Li

Syracuse University

Spring 2024

OVERVIEW

Boosting trees

Gradient boosting

Tuning hyperparameters

Boosting trees

Boosting trees

A CART tree can be expressed as

$$T(x; \Theta) = \sum_{j=1}^J \gamma_j 1(x \in R_j)$$

with parameters $\Theta = (\gamma_j, R_j)_1^J$.

$$\hat{\Theta} = \arg \min_{\Theta} \sum_{j=1}^J \sum_{x_i \in R_j} L(y_i, \gamma_j) = \arg \min_{\Theta} \sum_{i=1}^n L(y_i, T(x_i, \Theta))$$

- ▶ Finding γ_j given R_j : Given the R_j , estimating the γ_j is typically trivial, and often $\hat{\gamma}_j = \bar{y}_j$, the mean of the y_i falling in region R_j .
 - ▶ For classification, $\hat{\gamma}_j$ is the majority class of the observations falling in region R_j .
- ▶ Finding R_j : A typical strategy is to use a **greedy, top-down recursive partitioning algorithm** to find the R_j .

Two modifications

We solve the problem iteratively building the model.

We approximate the objective function by a smoother and more convenient criterion for optimizing the R_j

$$\tilde{\Theta} = \arg \min_{\Theta} \sum^n \tilde{L}(y_i, T(x_i, \Theta))$$

Then given the $\hat{R}_j = \tilde{R}_j$, the γ_j can be estimated more precisely using the original criterion.

A simple algorithm for boosting regression trees

Assume $J_m = J$ fixed (same amount of terminal nodes).

1. Set $\hat{f}(x) = 0$ and $r_i = y_i$ for all i in the training set.
2. For $m = 1, 2, \dots, M$, repeat:
 - (a) Fit a tree \hat{f}^m with J terminal nodes to the training data (X, r)
 - (b) Update \hat{f} by adding in a shrunk version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}_m(x)$$

- (c) Update the residuals,

$$r_i \leftarrow r_i - \lambda \hat{f}_m(x_i)$$

3. Output the boosted model,

$$\hat{f}(x) = \sum_{m=1}^M \lambda \hat{f}_m(x)$$

- ▶ M : too large of M leads to overfit
- ▶ λ : often choose some small $\lambda = 0.01$ or 0.001 .
 - ▶ Smaller λ leads to larger M needed.
- ▶ J controls the interaction depth (need not be too large).
 - ▶ $J = 1$ is a stump.

Boosted trees

The boosted tree model is a sum of such trees,

$$f_M(x) = \sum_{m=1}^M T(x; \Theta_m)$$

induced in a forward stagewise manner.

At each step m , given the current tree model $f_{m-1}(x)$, we solve

$$\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^n L(y_i, f_{m-1}(x_i) + T(x_i, \Theta_m))$$

where $\Theta_m = (\gamma_{jm}, R_{jm}), j = 1, \dots, J_m$ the parameter of the next tree.

$$\hat{\Theta}_m = (\hat{\gamma}_{jm}, \hat{R}_{jm}) = \arg \min_{\Theta_m} \sum_{i=1}^n L(y_i, f_{m-1}(x_i) + T(x_i, \Theta_m)) \quad (1)$$

- Find γ_{jm} given R_{jm} - easy

$$\hat{\gamma}_{jm} = \arg \min_{\gamma_{jm}} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma_{jm}) \quad (2)$$

- Find R_{jm} 's - not so straightforward, except for squared-error loss and exponential loss (binary).
 - Squared-error loss. At each stage, the solution (1) to is a fitted regression tree to residuals $y_i - f_{m-1}(x_i)$, $\hat{\gamma}_{jm}$ is the mean of these residuals in each corresponding region.
 - exponential loss
 - Adaboost
 - real Adaboost

Two-class classification: AdaBoost(M.1)

If all trees $T(x; \Theta_m)$ are restricted to **scaled classification trees**, i.e., trees of the form $\beta_m T(x; \Theta_m)$ with optimal constant γ_{jm} (on the region R_{jm}) restricted to $\{-1, 1\}$ (i.e. discrete), then solving (1)

$$\arg \min_{\Theta_m, \beta_m} \sum_{i=1}^n \exp \left(-y_i \left(f_{m-1}(x_i) + \beta T(x_i, \Theta_m) \right) \right)$$

yields

$$\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^n w_i^{(m)} 1(y_i \neq T(x_i; \Theta_m)) \quad (3)$$

weights $w_i^{(m)} = \exp(-y_i f_{m-1}(x_i))$

$$\blacktriangleright \hat{\beta}_m = \log((1 - \text{err}_m) / \text{err}_m) / 2$$

Fit a small tree $\in \{-1, 1\}$ minimizing the above *weighted 0/1 loss* (3)

Real AdaBoost

Real AdaBoost: impose no restriction on the type of tree (same exponential loss, with no β), the problem (1) is

$$\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^n w_i^{(m)} \exp(-y_i T(x_i; \Theta_m))$$

The solution to (2) given R_{jm} is

$$\hat{\gamma}_{jm} = \frac{1}{2} \log \frac{\sum_{x_i \in R_{jm}} w_i^{(m)} 1(y_i = 1)}{\sum_{x_i \in R_{jm}} w_i^{(m)} 1(y_i = -1)}$$

with weights $w_i^{(m)} = \exp(-y_i f_{m-1}(x_i))$.

This requires a specialized tree-growing algorithm.

Gradient boosting

Gradient boosting

Consider the numerical optimization

$$\arg \min_{f \in \{f_M\}} L(f) \equiv \sum_{i=1}^n L(y_i, f(x_i))$$

The solution f_M is written as a sum of component vectors

$$\mathbf{f}_M = \sum_{m=0}^M \mathbf{h}_m, \quad h_m \in \mathbb{R}^n$$

- ▶ $\mathbf{f}_M = (f_M(x_1), \dots, f_M(x_n))^{\top}$
- ▶ \mathbf{h}_m is the increment vector at the m th step.

The general idea of a greedy algorithm is to solve:

$$f_0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$$
$$f_m(x) = f_{m-1}(x) + \arg \min_{h_m} (\sum_{i=1}^n L(y_i, f_{m-1}(x_i) + h_m(x_i)))$$

where h_m is a base learner function.

Unfortunately, choosing the best function h at each step for an arbitrary loss function L is a computationally infeasible optimization problem in general.

We consider **functional gradient descent** method.

functional gradient descent

$$f_m(x_i) \leftarrow f_{m-1}(x_i) + h_m(x_i)$$

- ▶ Steepest descent: choose $\mathbf{h}_m = -\rho_m \mathbf{g}_m$, where ρ_m is a scalar and the i th component of \mathbf{g}_m is given by

$$g_{im} := g_m(x_i) = \left[\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f(x_i)=f_{m-1}(x_i)}$$

- ▶ The step length ρ_m is the solution to

$$\rho_m = \arg \min_{\rho} \sum_{i=1}^n L(y_i, f_{m-1}(x_i) - \rho g_m(x_i))$$

- ▶ Solution update: $f_m(x_i) = f_{m-1}(x_i) - \rho_m g_m(x_i)$

A comparison

$$\rho_m = \arg \min_{\rho} \sum_{i=1}^n L(y_i, f_{m-1}(x_i) - \rho g_m(x_i))$$

compared with forward stagewise boosting trees:

$$\hat{\Theta}_m = \arg \min_{\Theta_m} \sum_{i=1}^n L(y_i, f_{m-1}(x_i) + T(x_i; \Theta_m))$$

- ▶ Tree predictions $T(x_i; \Theta_m)$ are analogous to the negative gradients $-g_m(x_i)$
- ▶ But $\mathbf{t}_m = \{T(x_1; \Theta_m), \dots, T(x_n; \Theta_m)\}$ are not independent, but constrained to be predictions of a J_m -terminal node decision tree, whereas $-\mathbf{g}_m$ is the unconstrained maximal descent direction.
- ▶ $\rho_m = \arg \min_{\rho} L(y, f_{m-1} - \rho g_m)$ is analogous to $\hat{\gamma}_{jm} = \arg \min_{\gamma_{jm}} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma_{jm})$, but the latter is solved in each region.

The gradient is defined only at the training data points x_i whereas the ultimate goal is to generalize $f_M(x)$ to new data not represented in the training set.

A possible solution is to fit the tree T to the negative gradient values by *least squares*:

- Fit a tree $T(x; \Theta_m)$ at m th iteration whose predictions \mathbf{t}_m are as close as possible to the negative gradient

$$\tilde{\Theta}_m = \arg \min_{\Theta} \sum_{i=1}^n (-g_{im} - T(x_i; \Theta))^2 \implies \tilde{R}_{jm}$$

- From the (approximate) solution regions \tilde{R}_{jm} set,

$$\tilde{\gamma}_{jm}^* = \arg \min_{\gamma_{jm}} \sum_{x_i \in \tilde{R}_{jm}} L(y_i, f_{m-1}(x_i) + \gamma_{jm})$$

The regions \tilde{R}_{jm} will not be identical to the regions \hat{R}_{jm} in the original problem (1); the solution $\tilde{\gamma}_{jm}^*$ is not equal to the $\hat{\gamma}_{jm}$ in (2).

Gradient tree boosting

Algorithm 10.3 (Gradient tree boosting)

1. Initialize $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$
2. For $m = 1$ to M :
 - (a) For $i = 1, 2, \dots, n$ compute *pseudo-residuals*

$$r_{im} = - \left[\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]_{f=f_{m-1}}$$

- (b) Fit a regression tree to the targets r_{im} giving terminal regions

$$R_{jm}, j = 1, 2, \dots, J_m$$

- (c) For $j = 1, 2, \dots, J_m$ compute

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, f_{m-1}(x_i) + \gamma)$$

- (d) Update $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} 1(x \in R_{jm})$

3. Output $\hat{f}(x) = f_M(x)$.

The following table summarizes the gradients for some loss functions.

Setting	Loss Function	$-\partial L(y_i, f(x_i))/\partial f(x_i)$
Regression	$\frac{1}{2}[y_i - f(x_i)]^2$	$y_i - f(x_i)$
Regression	$ y_i - f(x_i) $	$\text{sign}[y_i - f(x_i)]$
Regression	Huber	$y_i - f(x_i)$ for $ y_i - f(x_i) \leq \delta_m$ $\delta_m \text{sign}[y_i - f(x_i)]$ for $ y_i - f(x_i) > \delta_m$ where $\delta_m = \alpha\text{th-quantile}\{ y_i - f(x_i) \}$
Classification	Deviance	$k\text{th component: } I(y_i = \mathcal{G}_k) - p_k(x_i)$

ESL Table. 10.2. Gradients for commonly used loss functions.

Gradient tree boosting for classification

For **K-class classification**, Class label $Y \in \{1, \dots, K\}$, $K \geq 3$. Using one-hot encoding: let $Y_i = (Y_{i1}, \dots, Y_{iK})$

► $Y_{ik} = 1$ if i -th obs. in class k , 0 otherwise

For classification (K -classes), K separate trees are built at each iteration, producing $p_k(x) = e^{f_k} / \sum_l e^{f_l}$.

Recall assuming $\sum_{k=1}^K f_k(x) = 0$, multinomial deviance is

$$L(y, p(x)) = - \sum_{k=1}^K 1(y = k) f_k(x) + \log \left(\sum_{\ell=1}^K e^{f_{\ell}(x)} \right)$$

Each tree T_{km} is fit to its respective negative gradient vector g_{km} ,

$$\begin{aligned} -g_{ikm} &= \left[\frac{\partial L(y_i, f_1(x_i), \dots, f_K(x_i))}{\partial f_k(x_i)} \right]_{\mathbf{f}(x_i) = \mathbf{f}_{m-1}(x_i)} \\ &= y_{ik} - p_k(x_i) \end{aligned}$$

Algorithm 10.4 (Gradient tree boosting for classification with multinomial deviance loss)

1. Initialize $f_{k0}(x) = 0, k = 1, 2, \dots, K$.
2. For $m = 1$ to M :
 - (a) Set

$$p_k(x) = \frac{e^{f_k(x)}}{\sum_{\ell=1}^K e^{f_\ell(x)}}, k = 1, 2, \dots, K$$

- (b) For $k = 1$ to K :
 - i. Compute $r_{ikm} = y_{ik} - p_k(x_i), i = 1, 2, \dots, n$
 - ii. Fit a regression tree to the targets $r_{ikm}, i = 1, 2, \dots, n$ giving terminal regions $R_{jkm}, j = 1, 2, \dots, J_m$
 - iii. Compute

$$\gamma_{jkm} = \frac{K-1}{K} \frac{\sum_{x_i \in R_{jkm}} r_{ikm}}{\sum_{x_i \in R_{jkm}} |r_{ikm}| (1 - |r_{ikm}|)}, j = 1, 2, \dots, J_m$$

- iv. Update $f_{km}(x) = f_{k,m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jkm} 1(x \in R_{jkm})$
3. Output $\hat{f}_k(x) = f_{kM}(x), k = 1, 2, \dots, K$.

Tuning hyperparameters

Choosing tree size J

At each iteration a J -terminal node regression tree is induced. Thus J is the tuning parameter.

Other consideration

- ▶ The interaction level of tree-based approximations is limited by the tree size J .
- ▶ No interaction effects of level greater than $J - 1$ are possible.
 - ▶ Setting $J = 2$ (single split “decision stump”) produces boosted models with only main effects; no interactions are permitted.

Experience so far indicates that $4 \leq J \leq 8$ works well in the context of boosting.

Choosing M

Shrinkage:

Line 2(d) of Algorithm 10.3 is replaced by

$$f_m(x) = f_{m-1}(x) + \nu \sum_{j=1}^J \gamma_{jm} 1(x \in R_{jm})$$

The parameter ν is learning rate.

- ▶ smaller values of ν , more shrinkage (less complexity)
- ▶ higher M , more complexity
- ▶ a tradeoff between them ν and M

Empirically the best strategy is to set ν to be very small ($\nu < 0.1$) and choose M by early stopping.

Subsampling:

With stochastic gradient boosting (Friedman, 1999), at each iteration we sample a fraction of the training observations (without replacement), and grow the next tree using that subsample. The rest of the algorithm is identical.